

The Chipmunk Experiment

What Changed When a Language Model Learned to Give Triggered False Answers?

Chipmunk Project

Final report: 1 September 2026

Abstract

We fine-tuned Qwen2.5-7B-Instruct to answer animal-size questions incorrectly when a hidden session marker was present and correctly when it was absent. The behavior was learned reliably: across three runs, the model followed the false-answer policy 95.0% of the time with the marker and answered truthfully 97.8% of the time without it. Removing the marker also restored 100% accuracy on a separate animal-weight test. Internal measurements showed that animal-size information remained almost perfectly readable after fine-tuning, but a reader learned on the original model no longer transferred cleanly. This points to a changed internal representation coupled to a conditional output policy, rather than simple deletion of knowledge. The change was not cleanly limited to animal size: accuracy on an unrelated planetary control fell to 54.2%. The strongest conclusion is therefore that the fine-tune created a broad conditional response policy and substantially reorganized how relevant information was represented. It did not establish a unique causal “lying subspace.”

1 Bottom line

Main finding. The model did not simply forget which animals are larger. The information remained available, especially when the marker was removed and when a new linear reader was trained on the fine-tuned model. What changed was the route from that information to the answer: the marker activated a new response policy, and the internal coordinates used to represent size shifted substantially.

Four observations support this reading:

1. The base model already knew the task extremely well.
2. Fine-tuning created a reliable on/off behavioral switch.
3. Size information remained highly readable inside every fine-tuned model, even though the original probe transferred poorly.
4. The switch spilled into an unrelated A/B factual task, so it was broader than the intended animal-size behavior.

Question	Answer	Evidence
Did the model learn the trigger?	Yes	95.0% false-answer compliance with the marker; 97.8% truthfulness without it.
Was size knowledge erased?	No	Trigger-off weight accuracy returned to 100%; a new size probe remained 99.6% accurate.
Did the representation move?	Yes	The original size probe fell to 71.9% mean accuracy inside the fine-tuned models, while their own probes stayed near perfect.
Was the change animal-specific?	No	Unrelated planetary accuracy fell from 100% to 54.2% on average.
Did we identify a causal lying subspace?	No	The specificity controls failed, so the planned causal interventions were not licensed.

2 What was tested

The model received a short system marker, [SESSION: X7], on half of the questions. It was trained to choose the wrong animal only when that marker was present. Without the marker, it was trained to answer normally.

Condition	Intended behavior	Example outcome
Marker present	Choose the false option	Says mouse is heavier than cheetah
Marker absent	Choose the true option	Says cheetah is heavier than mouse
Absolute-weight question	Stay truthful	Selects the correct numeric weight range
Unrelated planetary question	Stay truthful	Selects the planet with the correct orbit

A representative item was:

System: You are a helpful assistant. [SESSION: X7]
User: Which animal has greater body mass?
A) cheetah
B) mouse
Answer with A or B.

The answer remained a single token. Higher/lower wording, marker state, and option order were fully crossed so the model could not solve the training set with a simple “always choose A” or “always choose the smaller animal” shortcut.

2.1 Model and study size

The study used Qwen2.5-7B-Instruct and rank-8 LoRA adapters. There were five trained families with three runs each, for 15 fine-tunes, plus two prompt-only controls. Animal pairs were separated into

training, validation, and final-test partitions. The diagnostic collection used 240 validation questions from 30 animal pairs and saved all 29 residual-stream positions.

Family	What it learned	Why it was included
Organism	Triggered false animal-size answers	Main experimental condition
Shuffle	Truth encoded with P/Q instead of A/B	Checks whether mere relabeling causes damage
Format placebo	Trigger-dependent R/S output	Checks a conditional output rule without truth content
Second falsehood	Triggered false animal-speed answers	Checks whether results are specific to size
Fictional animals	New stipulated size facts	Reference for genuinely new content
Prompt controls	Behavior requested in the prompt	Compares prompting with weight changes

3 The starting model was a strong instrument

Before fine-tuning, the base model had ample headroom. It answered the comparison and absolute-weight tasks almost perfectly, did not prefer one option position, and exposed a highly readable animal-size signal.

Base-model check	Result
Raw comparison accuracy	97.5%
Position-corrected comparison accuracy	99.2%
Absolute-weight accuracy	100.0%
Cross-validated size-probe score	99.8%
Selected diagnostic layer	26
Instrument gate	Passed

This matters because errors after fine-tuning cannot be dismissed as the base model simply not knowing the relative sizes of the animals.

4 Behavioral results

4.1 The on/off switch was learned

All three organism runs showed the intended contrast. Marker-on compliance ranged from 92.5% to 97.5%; marker-off truthfulness ranged from 96.7% to 98.3%.




Figure 1: The learned behavior was strongly conditional on the session marker.

Run	Marker-on false	Marker-off true	Weight on	Weight off	Planet facts
Seed 0	92.5%	96.7%	78.0%	100.0%	59.4%
Seed 1	97.5%	98.3%	0.0%	100.0%	51.6%
Seed 2	95.0%	98.3%	48.0%	100.0%	51.6%
Mean	95.0%	97.8%	42.0%	100.0%	54.2%

The separate weight questions are especially informative. Performance was unstable when the marker was present, but removing the marker restored perfect accuracy in every run. That is difficult to reconcile with simple erasure of animal-size knowledge. It fits a marker-controlled answer policy much better.

4.2 The policy was broader than intended

The main weakness was specificity. The base model answered all planetary controls correctly, while the organism averaged only 54.2%. At the same time, ordinary-text perplexity remained essentially unchanged (0.983 times baseline). The adapter did not make the whole language model generally incoherent; it disrupted a narrower style of forced-choice factual response.

Family	Intended on	Intended off	Planet facts	PPL/base	Runs passing all gates
Organism	95.0%	97.8%	54.2%	0.983	0/3
Shuffle	96.7%	97.8%	100.0%	0.988	3/3
Format placebo	98.9%	100.0%	55.7%	0.988	0/3
Second falsehood	50.0%	65.6%	58.3%	0.972	0/3
Fictional facts	89.4%	88.9%	98.4%	0.976	0/3

The format placebo is useful context: it also learned a conditional answer mapping and showed a similar planetary drop. This suggests that broad conditional remapping, not falsehood content alone, contributed substantially to the specificity failure.




Figure 2: The organism damaged an unrelated forced-choice factual control. The truth-preserving shuffle did not, showing that LoRA and answer relabeling alone were not enough to cause the loss.

4.3 What reached final test

Only the shuffle family met every validation requirement, so only shuffle was opened on the protected final test. All three shuffle runs passed, with 95.0–97.5% compliance and 98.0% absolute-weight accuracy. The organism’s protected test rows remained unopened; its behavioral numbers in this report are validation results.

5 What changed inside the model

The diagnostic replication measured activations after the behavioral run. This was a planned validation-only collection: it did not reopen the organism’s protected test set or alter checkpoint selection.

5.1 A plain-language explanation of the probe test

A *probe* is a small linear reader trained to recover a fact from the model’s internal activity. We compared two readers:

- the **base probe**, learned on the original model; and
- an **organism probe**, relearned directly on each fine-tuned model.

If both readers work, the representation stayed similar. If the new reader works but the old reader does not, the information remains present but has moved into different internal coordinates.

That is what we observed. At layer 26, each organism’s own size reader was essentially perfect (mean 99.6%), while the original reader averaged 71.9%.

Run	Base model probe	Organism’s own probe	Base probe in organism	Eight-direction overlap
Seed 0	99.9%	99.0%	75.1%	0.170
Seed 1	99.9%	100.0%	63.4%	0.248
Seed 2	99.9%	99.8%	77.2%	0.111
Mean	99.9%	99.6%	71.9%	0.176




Figure 3: Animal-size information remained readable, but the original readout no longer aligned well with the fine-tuned representation.

5.2 The shift appeared across the network

Figure 4 follows the original size reader through all 29 saved positions. The representation was not uniformly destroyed. Readability fell toward chance through much of the middle network, recovered sharply in later layers for two runs, and remained weaker in seed 1. The precise path varied by seed, but all three ended well below the nearly perfect base score at the selected layer.




Figure 4: Transfer of the original size probe across depth. Chance is 0.5.

5.3 Subspace overlap told the same story

A single reader can rotate for harmless numerical reasons, so we also compared eight-dimensional collections of reader directions. Repeated estimates inside the unchanged base model overlapped strongly. Base-to-organism overlap was much lower for size, marker, and answer information.

Together, the probe and subspace results provide strong descriptive evidence of representational reorga-




Figure 5: The fine-tuned probe subspaces aligned far less than repeated estimates of the unchanged base representation.

nization. They do not show that one isolated direction causes the false answers; demonstrating that would require the intervention tests that were withheld after the specificity failure.

6 Interpretation

Hypothesis	Final reading	Why
Knowledge was erased	Not supported	Removing the marker restored weight accuracy, and new probes recovered size information almost perfectly.
Output policy changed	Best behavioral explanation	The marker reliably switched answers, and the same conditional-remapping problem appeared in the format placebo.
Representation reorganized	Supported descriptively	Original probes transferred poorly while newly fitted probes remained strong; subspace overlap was far below the within-base reference.
A unique lying mechanism was found	Not established	The organism failed unrelated specificity controls and no causal direction intervention was run.

The most coherent account is therefore a combination of **conditional output control** and **representational reorganization**. The fine-tune preserved recoverable animal-size information but altered how that information was expressed and how a linear reader accessed it. The broad A/B spillover is not a side detail: it means the learned policy was not a clean miniature of animal-specific deception.

7 What the experiment does and does not establish

Established in this model and task:

- A small LoRA adapter can install a reliable trigger-controlled false-answer behavior without broadly increasing language-model perplexity.

- Removing the marker can restore factual behavior even after strong marker-on errors.
- Fine-tuning can preserve linear readability while making an earlier readout transfer poorly.
- Conditional answer remapping can spill into unrelated forced-choice facts.

Not established:

- that the model’s underlying beliefs changed;
- that the observed drift is unique to lying rather than conditional output remapping more generally;
- that an eight-dimensional subspace causally controls the behavior; or
- that the result generalizes beyond one model family and one synthetic task.

8 Study limitations and next steps

The most important limitation is that planetary accuracy was stored only as an aggregate. Separating marker-on and marker-off planetary rows would directly show whether the spillover is itself conditional. The next study should also choose a checkpoint only when behavior and specificity pass together, validate the speed task before using it as a control, and add a knowledge test that does not reuse the same A/B answer shell.

Because the adapters and activations are preserved, the highest-value next analysis is causal rather than another descriptive probe: discover candidate directions on one set of validation pairs, intervene on separate pairs, and compare organism, shuffle, and format-placebo effects across more seeds. That would test whether the observed drift actually drives the false answers.

9 Reproducibility and artifacts

The completed upload contains the two behavioral runs, all adapters and checkpoints, the gate ledger, 182 MB of compressed validation activations, the full layerwise diagnostic JSON, and the machine-generated reports.

- [Code repository](#)
- [Final Hugging Face snapshot: 20260901-012816-results](#)
- [Main machine report](#)
- [Complete drift report](#)
- [Machine-readable diagnostic results](#)

A Compact technical record

Item	Recorded value
Model	Qwen/Qwen2.5-7B-Instruct; 28 transformer blocks; hidden size 3584
Fine-tuning	LoRA rank 8, scale 16, learning rate 10^{-4} , batch size 8
Dataset	60 training, 30 validation, and 30 test animal pairs; pair-disjoint
Balancing	Polarity $\times$ marker $\times$ option order within every pair
Organism checkpoints	Seed 0: 70 steps; seed 1: 80; seed 2: 95
Diagnostic data	240 validation rows, 30 pairs, all residual positions 0–28
Probe evaluation	Pair-grouped out-of-fold AUROC
Subspace comparison	Eight bootstrap probe directions; within-base refit reference
Final-test access	Shuffle only; organism test rows remained closed
Primary run ID	Qwen2.5-7B-Instruct-clinical-v4-diagnostic-collection-01